

DataSelf

Your AI learns how you work. You should be able to take that with you.

A local memory layer that turns real LLM sessions into portable working rules - without storing the conversation itself. github.com/nmndnkr/datasetself

01
primary user

0
raw chats retained

8
first-session memories

4/5
behaviors transferred

19/19
tests passing

01 / THE PRODUCT TURN

I built the counter first. It exposed the better product.

The hackathon version counted prompts, responses, copies, and corrections. That made activity visible, but activity was not the thing worth carrying between providers.

PRODUCT INSIGHT

The useful unit is not a prompt. It is a reusable working rule: how someone researches, what they verify, what they delegate, and what done means.

02 / USER + JOB

One user. One painful reset.

First user
An individual whose LLM becomes more useful as it learns how they operate.

Job to be done
When I change LLMs, bring how I work so I do not have to teach the new one from zero.

Why not teams first?
Permissions and shared truth widen the system before the core transfer is proven.

03 / MVP LOOP

Five steps. One portable artifact.

- 1

Start - Toggle the macOS menu app.
- 2

Work - Use the active LLM normally.
- 3

Extract - Request a strict memory patch.
- 4

Validate - Reject unsafe or invalid output.
- 5

Move - Export datasetself-context.md.

04 / FIRST THREE PRODUCT DECISIONS

CHOSE Behavioral rules	OVER Raw transcripts	WHY Portable and legible without retaining the most sensitive material.	COST Less nuance.
CHOSE Active LLM extraction	OVER A second model	WHY The provider already has the session, so no second copy is needed.	COST UI dependency.
CHOSE Local Markdown	OVER Cloud account	WHY The user can inspect, edit, move, or delete every memory.	COST No sync.

04 / PRODUCT DECISIONS CONTINUED

CHOSE Automatic + provisional	OVER Approve every item	WHY A menu-bar tool cannot become another inbox; users can still edit or undo.	COST Weak memory may appear.
CHOSE One context file	OVER Provider connectors	WHY Markdown works everywhere and made migration testable in one week.	COST Manual first step.

05 / CONTROLLED MIGRATION TEST

I tested behavior, not whether the file opened.

Two fresh anonymous ChatGPT sessions received the same vague product request. One had no memory; one received DataSelf's context.

WITHOUT DATASELF
Selected a direction, proposed a broad multi-week build, and asked several questions at once.
WITH DATASELF
Made no product assumption, asked one focused question, kept scope narrow, and reserved execution.

Result: 4 of 5 behaviors transferred without correction. The completion-handoff behavior was not exercised because the test stopped at definition.

06 / SUCCESS METRICS

Success is less re-explanation.

NORTH STAR	Less workflow re-explanation	Outcome
KNOWN PREFERENCES	At least 4 of 5	Passed: 4/5
RAW CHAT RETAINED	0 bytes	Passed
EVIDENCE LINKS	100%	Passed
FIRST UPDATE	Under 10 minutes	Unproven
CORRECTIVE MESSAGES	Below baseline	Next test

07 / WHAT IS PROVEN - AND WHAT IS NOT

Built evidence is not market evidence.

The prototype proves that behavioral memory can transfer. It does not prove that enough people feel the reset, trust automatic extraction, or keep using the system.

Next experiment

Five heavy LLM users, three sessions each. Measure setup time, memory edit/deletion rate, corrective messages with and without context, and repeat migration.

Risks to watch

Provider UI churn / confident but wrong memory / context growth becoming noise / convenience mattering more than portability.